

Details

Result

Result

Time

Cycles

Regs

GPU

SM Frequency

OC

Process

Save as PDF

Current

m.m.0

538 - shared_mem_m...

31.20 msec

46,802,969

27

0 - NVIDIA GeForce RTX 3070

1.50 cycle/nsec

8.6

[16620] HW3 exe

100x100_global

m.m.0

535 - global_mem_m...

1.86 msec

2,791,525

40

0 - NVIDIA GeForce RTX 3070

1.50 cycle/nsec

8.6

[15504] HW3 exe

100x100_shared

m.m.0

538 - shared_mem_m...

2.14 msec

3,215,351

27

0 - NVIDIA GeForce RTX 3070

1.50 cycle/nsec

8.6

[15504] HW3 exe

1000x1000_global

1000x1000_shared

2000x3000_global

2000x3000_shared

GPU Speed Of Light Throughput

All

High-level overview of the throughput for compute and memory resources of the GPU. For each unit, the throughput reports the achieved percentage of utilization with respect to the theoretical maximum. Breakdowns show the throughput for each individual sub-metric of Compute and Memory to clearly identify the highest contributor. High-level overview of the utilization for compute and memory resources of the GPU presented as a profile chart.

Metric	Value	Delta	Target
Memory Throughput [%]	98.76	(+3.47%, z=+0.47)	100.00
Memory Throughput [GB/s]	98.76	(+3.47%, z=+0.47)	100.00
L1/TEX Cache Throughput [%]	98.81	(-0.09%, z=+0.19)	100.00
L2 Cache Throughput [%]	17.29	(-2.75%, z=+0.27)	20.00
DRAM Throughput [%]	29.81	(+112.82%, z=+1.05)	27.30

High Throughput

The kernel is utilized greater than 80.0% of the available compute or memory performance of the device. To further improve performance, work will likely need to be shifted from the most utilized to another unit. Start by analyzing the workload in the [Kernel Profiling Guide](#) section.

Routine Analysis

The ratio of peak float (fp32) to double (fp64) performance on this device is 64:1. The kernel achieved 5% of this device's fp32 peak performance and 0% of its fp64 peak performance. See the [Kernel Profiling Guide](#) for more details on routine analysis.

GPU Throughput

Compute (SM) [%]

Memory [%]

Speed of Light (SOL) [%]

Compute Throughput Breakdown

Memory Throughput Breakdown

Floating Point Operations Routine

Performance (1/DP/12) (1/1E+12)

Arithmetic Intensity (FLOP/byte)

Compute Workload Analysis

All

Detailed analysis of the compute resources of the streaming multiprocessors (SM), including the achieved instructions per clock (IPC) and the utilization of each available pipeline. Pipelines with very high utilization might limit the overall performance.

Metric	Value	Delta	Target
Executed Ipc Elapsed [inst/cycle]	1.40	(+6.00%, z=+0.86)	1.32
Executed Ipc Active [inst/cycle]	1.40	(+12.42%, z=+0.88)	1.24
Issued Ipc Active [inst/cycle]	1.40	(+12.42%, z=+0.88)	1.24

Balanced

ALU is the highest-utilized pipeline (25.4%) based on active cycles, taking into account the rates of its different instructions. It executes integer and logic operations. It is well-utilized, but should not be a bottleneck.

Pipe Utilization (% of active cycles)

Utilization [%]

Pipe Utilization (% of peak instructions executed)

Utilization [%]

Memory Workload Analysis

All

Detailed analysis of the memory resources of the GPU. Memory can become a limiting factor for the overall kernel performance when fully utilizing the involved hardware units (Mem Bus), causing the fail to achieve communication bandwidth between those units (Max Bandwidth), or by reaching the maximum throughput of issuing memory instructions (Mem Pipes Busy).

Metric	Value	Delta	Target
Memory Throughput [GB/s/second]	129.62	(+112.76%, z=+1.05)	115.00
L1/TEX Hit Rate [%]	11.00	(-77.23%, z=+0.82)	50.00
L2 Hit Rate [%]	48.49	(-36.02%, z=+1.15)	85.00
L2 Compression Success Rate [%]	0	(+0.00%, z=+0.00)	100.00

L1/TEX Global Load Access Pattern

The memory access pattern for global loads in L1/TEX might not be optimal. On average, this kernel accesses 4.0 bytes per thread per memory request, but the address pattern, possibly caused by the stride between threads, results in 4.1 sectors per request, or 4.1*32 = 132.4 bytes of cache data transfers per request. The optimal thread address pattern for 4.0 byte accesses would result in 4.0*32 = 128.0 bytes of cache data transfers per request, to maximize L1/TEX cache performance. Check the [Source Counters](#) section for uncoalesced global loads.

L2 Load Access Pattern

The memory access pattern for loads from L1/TEX to L2 is not optimal. The granularity of an L1/TEX request to L2 is a 128-byte cache line. That is 4 consecutive 32-byte sectors per L2 request. However, this kernel only accesses an average of 1.6 sectors out of the possible 4 sectors per memory request. Check the [Source Counters](#) section for uncoalesced loads and try to minimize how many cache lines need to be accessed per memory request.

L2 Store Access Pattern

The memory access pattern for stores from L1/TEX to L2 is not optimal. The granularity of an L1/TEX request to L2 is a 128-byte cache line. That is 4 consecutive 32-byte sectors per L2 request. However, this kernel only accesses an average of 1.6 sectors out of the possible 4 sectors per memory request. Check the [Source Counters](#) section for uncoalesced stores and try to minimize how many cache lines need to be accessed per memory request.

Memory Chart

Show As: Transfer Size

Global

Local

Texture

Surface

Load Global Store Shared

Shared

L1/TEX Cache

L2 Cache

System Memory

Device Memory

DRAM

Shared Memory

Shared Memory

Instructions

Requests

Wavefronts

% Peak

Bank Conflicts

L1/TEX Cache

Instructions

Requests

Wavefronts

% Peak

Sectors

Sectors/Req

L2 Cache Eviction Policies

First

Hit Rate

Last

Hit Rate

Normal

Hit Rate

Bytes

Throughput

Schedule Statistics

Summary of the activity of the schedulers issuing instructions. Each scheduler maintains a pool of warps that it can issue instructions for. The upper bound of warps in the pool (Theoretical Warps) is limited by the launch configuration. On every cycle each scheduler checks the state of the allocated warps in the pool (Active Warps). Active warps that are not stalled (Theoretical Warps) are ready to issue their next instruction. From the set of eligible warps the scheduler selects a single warp from which to issue one or more instructions (Issued Warp). On cycles with no eligible warps, the issue slot is skipped and no instruction is issued. Having many skipped issue slots indicates poor latency issue. On cycles with no eligible warps, the issue slot is skipped and no instruction is issued. Having many skipped issue slots indicates poor latency issue. To increase the number of eligible warps, avoid possible load imbalances due to highly different execution durations per warp. Reducing stalls indicated on the [Warp State Statistics](#) and [Source Counters](#) sections can help too.

Metric	Value	Delta	Target
Active Warps Per Scheduler [warp]	11.94	(+1.34%, z=+1.15)	10.00
Eligible Warps Per Scheduler [warp]	1.48	(-11.73%, z=+0.20)	1.00
Issued Warp Per Scheduler	0.35	(+12.42%, z=+0.88)	0.30

Issue Slot Utilization

Every scheduler is capable of issuing one instruction per cycle, but for this kernel each scheduler only issues an instruction every 2.9 cycles. This might leave hardware resources underutilized and may lead to less optimal performance. Out of the maximum of 12 warps per scheduler, the kernel allocates an average of 11.94 active warps per scheduler, but only an average of 1.48 warps were eligible per cycle. Eligible warps are the subset of active warps that are ready to issue their next instruction. Every cycle with no eligible warps results in no instruction being issued and the issue slot remains unused. To increase the number of eligible warps, avoid possible load imbalances due to highly different execution durations per warp. Reducing stalls indicated on the [Warp State Statistics](#) and [Source Counters](#) sections can help too.

Warp State Statistics

GPU Maximum Warps Per Scheduler

Theoretical Warps Per Scheduler

Active Warps Per Scheduler

Eligible Warps Per Scheduler

Issued Warp Per Scheduler

Warp State Statistics

Analysis of the states in which all warps spent cycles during the kernel execution. The warp states describe a warps readiness or inability to issue its next instruction. The warp cycles per instruction define the latency between two consecutive instructions. The higher the value, the more warp parallelism is required to hide this latency. For each warp state, the chart shows the average number of cycles spent in that state per issued instruction. Stalls are not always impacting the overall performance nor are they completely avoidable. Only focus on stall reasons if the schedulers fail to issue every cycle. When executing a kernel with mixed library and user code, these metrics show the combined values.

Metric	Value	Delta	Target
Warp Cycles Per Issued Instruction [cycle]	34.09	(-11.01%, z=+0.79)	38.00
Warp Cycles Per Executed Instruction [cycle]	34.09	(-11.01%, z=+0.79)	38.00

Warp State Statistics

On average, each warp of this kernel spends 18.6 cycles being stalled waiting for an MIO instruction queue to be not full. This represents about 54.7% of the total average of 34.1 cycles between issuing two instructions. This stall reason is high in cases of utilization of the MIO pipelines, which include special math instructions, dynamic branches, as well as shared memory instructions. When caused by shared memory accesses, trying to use fewer but wider loads can reduce pipeline pressure.

Warp Stall

Check the [Source Counters](#) section for the top stall locations in your source based on sampling data. The [Kernel Profiling Guide](#) provides more details on each stall reason.

Warp State (All Cycles)

Stall MIO Throttle

Stall Barrier

Stall Long Scoreboard

Stall Not Selected

Stall Wait

Selected

Stall Short Scoreboard

Stall Dispatch Stall

Stall Math Pipe Throttle

Stall No Instruction

Stall Branch Resolving

Stall Drain

Stall LG Throttle

Stall IMC Miss

Stall Tex Throttle

Stall Misc

Stall Member

Stall Sleeping

Instruction Statistics

Statistics of the executed low-level assembly instructions (SASS). The instruction mix provides insight into the types and frequency of the executed instructions. A narrow mix of instruction types implies a dependency on one instruction pipeline, while others remain unused. Using multiple pipelines allows hiding latencies and enables parallel execution. Note that Instructions/Opcode and Executed Instructions are measured differently and can diverge if cycles are spent in system calls.

Metric	Value	Delta	Target
Executed Instructions [inst]	3,015,896,000	(+214.91%, z=+1.37)	1,000,000,000
Issued Instructions [inst]	3,015,916,076	(+214.91%, z=+1.37)	1,000,000,000

Executed Instruction Mix

LD

LDL

LDL2

LDL3

LDL4

LDL5

LDL6

LDL7

LDL8

LDL9

LDL10

LDL11

LDL12

LDL13

LDL14

LDL15

LDL16

LDL17

LDL18

LDL19

LDL20

LDL21

LDL22

LDL23

LDL24

LDL25

LDL26

LDL27

LDL28

LDL29

LDL30

LDL31

LDL32

LDL33

LDL34

LDL35

LDL36

LDL37

LDL38

LDL39

LDL40

LDL41

LDL42

LDL43

LDL44

LDL45

LDL46

LDL47

LDL48

LDL49

LDL50

LDL51

LDL52

LDL53

LDL54

LDL55

LDL56

LDL57

LDL58

LDL59

LDL60

LDL61

LDL62

LDL63

LDL64

LDL65

LDL66

LDL67

LDL68

LDL69

LDL70

LDL71

LDL72

LDL73

LDL74

LDL75

LDL76

LDL77

LDL78

LDL79

LDL80

LDL81

LDL82

LDL83

LDL84

LDL85

LDL86

LDL87

LDL88

LDL89

LDL90

LDL91

LDL92

LDL93

LDL94

LDL95

LDL96

LDL97

LDL98

LDL99

LDL100

LDL101

LDL102

LDL103

LDL104

LDL105

LDL106

LDL107

LDL108

LDL109

LDL110

LDL111

LDL112

LDL113

LDL114

LDL115

LDL116

LDL117

LDL118

LDL119

LDL120

LDL121

LDL122

LDL123

LDL124

LDL125

LDL126

LDL127

LDL128

LDL129

LDL130

LDL131

LDL132

LDL133

LDL134

LDL135

LDL136

LDL137

LDL138

LDL139

LDL140

LDL141

LDL142

LDL143

LDL144

LDL145

LDL146

LDL147

LDL148

LDL149

LDL150

LDL151

LDL152

LDL153

LDL154

LDL155

LDL156

LDL157

LDL158

LDL159

LDL160

LDL161

LDL162

LDL163

LDL164

LDL165

LDL166

LDL167

LDL168

LDL169

LDL170

LDL171

LDL172

LDL173

LDL174

LDL175

LDL176

LDL177

LDL178

LDL179

LDL180

LDL181

LDL182

LDL183

LDL184

LDL185

LDL186

LDL187

LDL188

LDL189

LDL190

LDL191

LDL192

LDL193

LDL194

LDL195

LDL196

LDL197

LDL198

LDL199

LDL200

LDL201

LDL202

LDL203

LDL204

LDL205

LDL206

LDL207

LDL208

LDL209

LDL210

LDL211

LDL212

LDL213

LDL214

LDL215

LDL216

LDL217

LDL218

LDL219

LDL220

LDL221

LDL222

LDL223

LDL224

LDL225

LDL226

LDL227

LDL228

LDL229

LDL230

LDL231

LDL232

LDL233

LDL234

LDL235

LDL236

LDL237

LDL238

LDL239

LDL240

LDL241

LDL242

LDL243

LDL244

LDL245

LDL246

LDL247

LDL248

LDL249

LDL250

LDL251

LDL252

LDL253

LDL254

LDL255

LDL256

LDL257

LDL258

LDL259

LDL260

LDL261

LDL262

LDL263

LDL264

LDL265

LDL266

LDL267

LDL268

LDL269

LDL270

LDL271

LDL272

LDL273

LDL274

LDL275

LDL276

LDL277

LDL278

LDL279

LDL280

LDL281

LDL282

LDL283

LDL284

LDL285

LDL286

LDL287

LDL288

LDL289

LDL290

LDL291

LDL292

LDL293

LDL294

LDL295

LDL296

LDL297

LDL298

LDL299

LDL300

LDL301

LDL302

LDL303

LDL304

LDL305

LDL306

LDL307

LDL308

LDL309

LDL310

LDL311

LDL312

LDL313

LDL314

LDL315

LDL316

LDL317

LDL318

LDL319

LDL320

LDL321

LDL322

LDL323

LDL324

LDL325

LDL326

LDL327

LDL328

LDL329

LDL330

LDL331

LDL332

LDL333

LDL334

LDL335

LDL336

LDL337

LDL338

LDL339

LDL340

LDL341

LDL342

LDL343

LDL344

LDL345

LDL346

LDL347

LDL348

LDL349

LDL350

LDL351

LDL352

LDL353

LDL354

LDL355

LDL356

LDL357

LDL358

LDL359

LDL360

LDL361

LDL362

LDL363

LDL364

LDL365

LDL366

LDL367

LDL368

LDL369

LDL370

LDL371

LDL372

LDL373

LDL374

LDL375

LDL376

LDL377

LDL378

LDL379

LDL380

LDL381

LDL382

LDL383

LDL384

LDL385

LDL386

LDL387

LDL388

LDL389

LDL390

LDL391

LDL392

LDL393

LDL394

LDL395

LDL396

LDL397

LDL398

LDL399

LDL400

LDL401

LDL402

LDL403

LDL404

LDL405

LDL406

LDL407

LDL408

LDL409

LDL410

LDL411

LDL412

LDL413

LDL414

LDL415

LDL416

LDL417

LDL418

LDL419

LDL420

LDL421

LDL422

LDL423

LDL424

LDL425

LDL426

LDL427

LDL428

LDL429

LDL430

LDL431

LDL432

LDL433

LDL434

LDL435

LDL436

LDL437

LDL438

LDL439

LDL440

LDL441

LDL442

LDL443

LDL444

LDL445

LDL446

LDL447

LDL448

LDL449

LDL450

LDL451

LDL452

LDL453

LDL454

LDL455

LDL456

LDL457

LDL458

LDL459

LDL460

LDL461

LDL462

LDL463

LDL464

LDL465

LDL466

LDL467

LDL468

LDL469

LDL470

LDL471

LDL472

LDL473

LDL474

LDL475

LDL476

LDL477

LDL478

LDL479

LDL480

LDL481

LDL482

LDL483

LDL484

LDL485

LDL486

LDL487

LDL488

LDL489

LDL490

LDL491

LDL492

LDL493

LDL494

LDL495

LDL496

LDL497

LDL498

LDL499

LDL500

LDL501

LDL502

LDL503

LDL504

LDL505

LDL506

LDL507

LDL508

LDL509

LDL510

LDL511

LDL512

LDL513

LDL514

LDL515

LDL516

LDL517

LDL518

LDL519

LDL520

LDL521

LDL522

LDL523

LDL524

LDL525

LDL526

LDL527

LDL528

LDL529

LDL530

LDL531

LDL532

LDL533

LDL534

LDL535

LDL536

LDL537

LDL538

LDL539

LDL540

LDL541

LDL542

LDL543

LDL544

LDL545

LDL546

LDL547

LDL548

LDL549

LDL550

LDL551

LDL552

LDL553

LDL554

LDL555

LDL556

LDL557

LDL558

LDL559

LDL560

LDL561

LDL562

LDL563

LDL564

LDL565

LDL566

LDL567

LDL568

LDL569

LDL570

LDL571

LDL572

LDL573

LDL574

LDL575

LDL576

LDL577

LDL578

LDL579

LDL580

LDL581

LDL582

LDL583

LDL584

LDL585

LDL586

LDL587

LDL588

LDL589

LDL590

LDL591

LDL592

LDL593

LDL594

LDL595

LDL596

LDL597

LDL598

LDL599

LDL600

LDL601

LDL602

LDL603

LDL604

LDL605

LDL606

LDL607

LDL608

LDL609

LDL610

LDL611

LDL612

LDL613

LDL614

LDL615

LDL616

LDL617

LDL618

LDL619

LDL620

LDL621

LDL622

LDL623

LDL624

LDL625

LDL626

LDL627

LDL628

LDL629

LDL630

LDL631

LDL632

LDL633

LDL634

LDL635

LDL636

LDL637

LDL638

LDL639

LDL640

LDL641

LDL642

LDL643

LDL644

LDL645

LDL646

LDL647

LDL648

LDL649

LDL650

LDL651

LDL652

LDL653

LDL654

LDL655

LDL656

LDL657

LDL658

LDL659

LDL660

LDL661

LDL662

LDL663

LDL664

LDL665

LDL666

LDL667

LDL668

LDL669

LDL670

LDL671

LDL672

LDL673

LDL674

LDL675

LDL676

LDL677

LDL678

LDL679

LDL680

LDL681

LDL682

LDL683

LDL684

LDL685

LDL686

LDL687

LDL688

LDL689

LDL690

LDL691

LDL692

LDL693

LDL694

LDL695

LDL696

LDL697

LDL698

LDL699

LDL700

LDL701

LDL702

LDL703

LDL704

LDL705

LDL706

LDL707

LDL708

LDL709

LDL710

LDL711

LDL712

LDL713

LDL714

LDL715

LDL716

LDL717

LDL718

LDL719

LDL720

LDL721

LDL722

LDL723

LDL724

LDL725

LDL726

LDL727

LDL728

LDL729

LDL730

LDL731

LDL732

LDL733

LDL734

LDL735

LDL736

LDL737

LDL738

LDL739

LDL740

LDL741

LDL742

LDL743

LDL744

LDL745

LDL746

LDL747

LDL748

LDL749

LDL750

LDL751

LDL752

LDL753

LDL754

LDL755

LDL756

LDL757

LDL758

LDL759

LDL760

LDL761

LDL762

LDL763

LDL764

LDL765

LDL766

LDL767

LDL768

LDL769

LDL770

LDL771

LDL772

LDL773

LDL774

LDL775

LDL776

LDL777

LDL778

LDL779

LDL780

LDL781

LDL782

LDL783

LDL784

LDL785

LDL786

LDL787

LDL788

LDL789

LDL790

LDL791

LDL792

LDL793

LDL794

LDL795

LDL796

LDL797

LDL798

LDL799

LDL800

LDL801

LDL802

LDL803

LDL804

LDL805

LDL806

LDL807

LDL808

LDL809

LDL810

LDL811

LDL812

LDL813

LDL814

LDL815

LDL816

LDL817

LDL818

LDL819